

Deep Learning Project

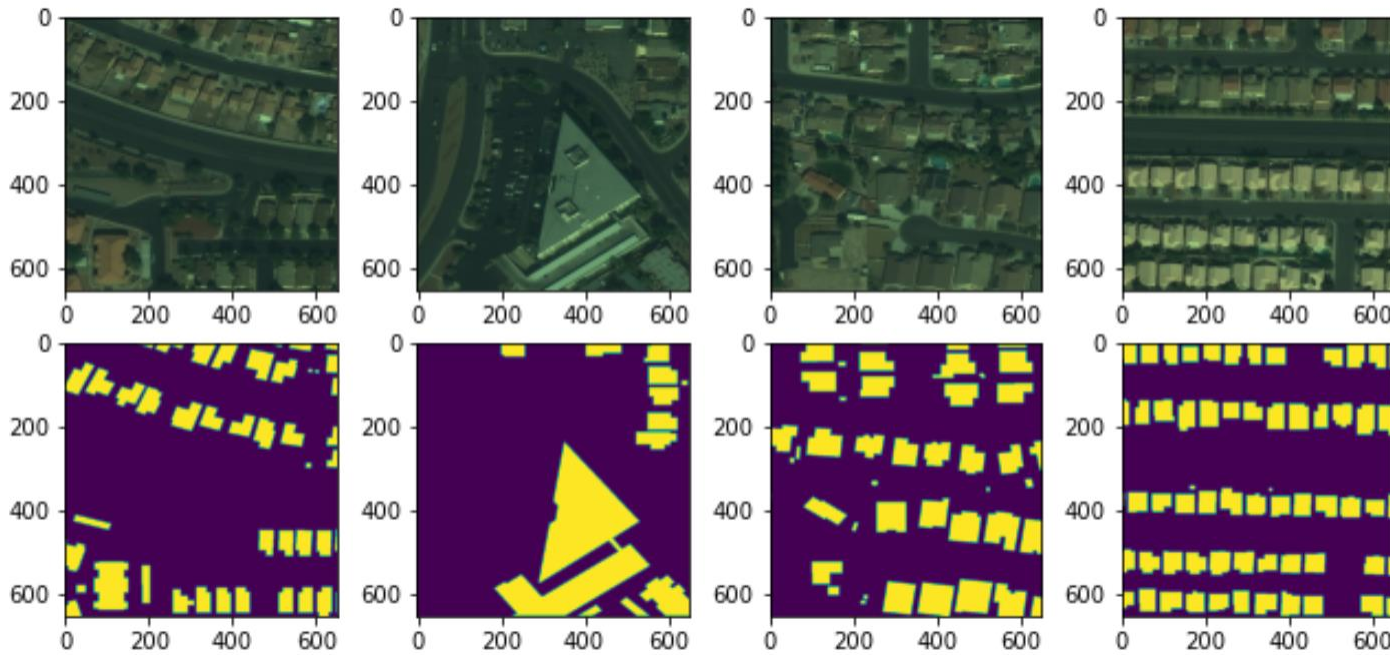
Harold Gamarro

Joe Brian Malubay

Said Mejia

Juan Pablo Montoya

Spacenet Data processing for building detection



Input Image:

- 3 band geotiff
- 650 x 650 x 3
- ~ 3000 images

Image and mask generator

Input mask:

- Geojson
- 650 x 650 x 3

Output data:

- Processed NumPy arrays in batches of 32 images
 - Normalized data
 - Rescaled to 512 x 512
-
- 91 files of 32x512x512x3 for images
 - 91 files of 32x512x512x1 for imagmaskses

Spacenet Data processing for building detection

Current Progress:

- Finished source code to convert images to jpegs for model input
- Tested out the framework to load image and mask into a Unet model that successfully trained model over 5 epochs
- Computation served as a big challenge resulting in OOM errors
 - Figured out how to use research compute cluster to handle the training
 - 32 cores with about 60 GB of memory
 - Can submit jobs and test various models

```
Epoch 1/5
91/91 [=====] - 4368s 48s/step - loss: 1.0767 - accuracy: 0.8102 - val_loss: 0.5613 - val_accuracy: 0.7996

Epoch 00001: val_loss improved from inf to 0.56129, saving model to /home/hgamarro/DeepLearning/HG_space/notebooks/random/model_2_che
Epoch 2/5
91/91 [=====] - 4356s 48s/step - loss: 0.2650 - accuracy: 0.8944 - val_loss: 0.8278 - val_accuracy: 0.7996

Epoch 00002: val_loss did not improve from 0.56129
Epoch 3/5
91/91 [=====] - 4354s 48s/step - loss: 0.2079 - accuracy: 0.9163 - val_loss: 0.9882 - val_accuracy: 0.7996

Epoch 00003: val_loss did not improve from 0.56129
Epoch 4/5
91/91 [=====] - 4334s 48s/step - loss: 0.1832 - accuracy: 0.9258 - val_loss: 0.8948 - val_accuracy: 0.7996

Epoch 00004: val_loss did not improve from 0.56129
Epoch 5/5
91/91 [=====] - 4348s 48s/step - loss: 0.1710 - accuracy: 0.9314 - val_loss: 1.1548 - val_accuracy: 0.7996

Epoch 00005: val_loss did not improve from 0.56129

Time Taken for testing: 6:02:45.082028
```

```
top - 16:03:47 up 263 days, 8:51, 0 users, load average: 12.74, 13.69, 13.07
Tasks: 418 total, 2 running, 416 sleeping, 0 stopped, 0 zombie
%Cpu0  : 47.8 us, 4.7 sy, 0.0 ni, 47.5 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu1  : 45.7 us, 7.9 sy, 0.0 ni, 46.4 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu2  : 43.9 us, 7.0 sy, 0.0 ni, 49.2 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu3  : 32.9 us, 21.9 sy, 0.0 ni, 45.2 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu4  : 45.2 us, 5.6 sy, 0.0 ni, 49.2 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu5  : 45.0 us, 5.6 sy, 0.0 ni, 49.3 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu6  : 47.5 us, 27.6 sy, 0.0 ni, 24.9 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu7  : 47.0 us, 4.0 sy, 0.0 ni, 49.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu8  : 46.7 us, 3.0 sy, 0.0 ni, 50.3 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu9  : 45.9 us, 4.0 sy, 0.0 ni, 50.2 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu10 : 46.4 us, 9.3 sy, 0.0 ni, 59.9 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu11 : 28.4 us, 7.6 sy, 0.0 ni, 64.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu12 : 28.5 us, 7.6 sy, 0.0 ni, 63.9 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu13 : 28.1 us, 7.9 sy, 0.0 ni, 63.9 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu14 : 29.2 us, 7.0 sy, 0.0 ni, 63.8 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu15 : 28.2 us, 7.6 sy, 0.0 ni, 64.1 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu16 : 27.5 us, 9.6 sy, 0.0 ni, 62.9 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu17 : 27.6 us, 10.0 sy, 0.0 ni, 62.5 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu18 : 25.5 us, 10.3 sy, 0.0 ni, 64.2 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu19 : 25.8 us, 9.9 sy, 0.0 ni, 64.2 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu20 : 25.7 us, 10.6 sy, 0.0 ni, 63.7 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu21 : 26.8 us, 11.3 sy, 0.0 ni, 61.9 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu22 : 25.2 us, 11.0 sy, 0.0 ni, 63.8 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu23 : 26.9 us, 9.0 sy, 0.0 ni, 64.1 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu24 : 39.9 us, 21.6 sy, 0.0 ni, 38.5 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu25 : 27.2 us, 9.3 sy, 0.0 ni, 63.6 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu26 : 28.4 us, 9.9 sy, 0.0 ni, 61.4 id, 0.3 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu27 : 26.4 us, 9.9 sy, 0.0 ni, 63.7 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu28 : 39.4 us, 18.9 sy, 0.0 ni, 41.7 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu29 : 27.2 us, 8.9 sy, 0.0 ni, 63.9 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu30 : 29.7 us, 16.2 sy, 0.0 ni, 54.1 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu31 : 52.5 us, 15.5 sy, 0.0 ni, 32.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 13184244+total, 3810656 free, 37913840 used, 90117952 buff/cache
KiB Swap: 24001104 total, 23641888 free, 359216 used, 92572672 avail Mem
```

- Each Epoch takes about 1 hour 20 min
- Uses all 32 cores
- Memory peaks 40 GB in above

```

In [4]: 1 def npy_generator(path ,npv_len=2):
        2     #for i in np.arange(npy_len):
        3     for fname in sorted(os.listdir(path)):
        4         if fname.endswith(".npy"):
        5             #print(fname)
        6             yield np.load(path+"/"+fname)
        7
        8 def masks_generator(path ,npv_len=2):
        9     #for i in np.arange(npy_len):
        10    for fname in sorted(os.listdir(path)):
        11        if fname.endswith(".npy"):
        12            #print(fname)
        13            yield np.load(path+"/"+fname)

```

```

In [5]: 1 target_img_paths

```

```

Out[5]: '/mnt/hgfs/VMsharedFolder/git/misc/masks'

```

```

In [6]: 1 def xy_generator(targ_data='/mnt/hgfs/VMsharedFolder/git/misc/npy'
        2                 ,targ_masks='/mnt/hgfs/VMsharedFolder/git/misc/masks'):
        3     for item1 ,item2 in zip(npy_generator(path=input_imgs_path)
        4                           ,npv_generator(path=target_img_paths)):
        5         yield(item1 ,item2)
        6         #print(item1.shape ,"|",item2.shape ,"/")
        7         #print(item1.shape ,"|",item2.shape ,"/")

```

```

In [7]: 1 plt.hist(np.load(item1)[0] + 0.1 , bins=100)

```


```
save_model_weights_hdf5(model, weight_path, overwrite = TRUE)
```

In [15]:

```
1 #copies all layers of model
2 activation_layers = [layer.output for layer in unet.layers[:]]
3 # Creates a model that will return these outputs, given the model input
4 activation_model = models.Model(inputs=unet.input, outputs=activation_layers)
5 # Returns a list of five Numpy arrays: one array per layer activation
6 activation_model.summary()
```

conv2d_transpose_7 (Conv2DTrans	(None, 256, 256, 32) 9248	activation_14[0][0]
batch_normalization_14 (BatchNo	(None, 256, 256, 32) 128	conv2d_transpose_7[0][0]
up_sampling2d_7 (UpSampling2D)	(None, 512, 512, 64) 0	add_5[0][0]
up_sampling2d_6 (UpSampling2D)	(None, 512, 512, 32) 0	batch_normalization_14[0][0]
conv2d_7 (Conv2D)	(None, 512, 512, 32) 2080	up_sampling2d_7[0][0]
add_6 (Add)	(None, 512, 512, 32) 0	up_sampling2d_6[0][0] conv2d_7[0][0]
conv2d_8 (Conv2D)	(None, 512, 512, 1) 289	add_6[0][0]

=====
Total params: 2,058,401
Trainable params: 2,054,625
Non-trainable params: 3.776



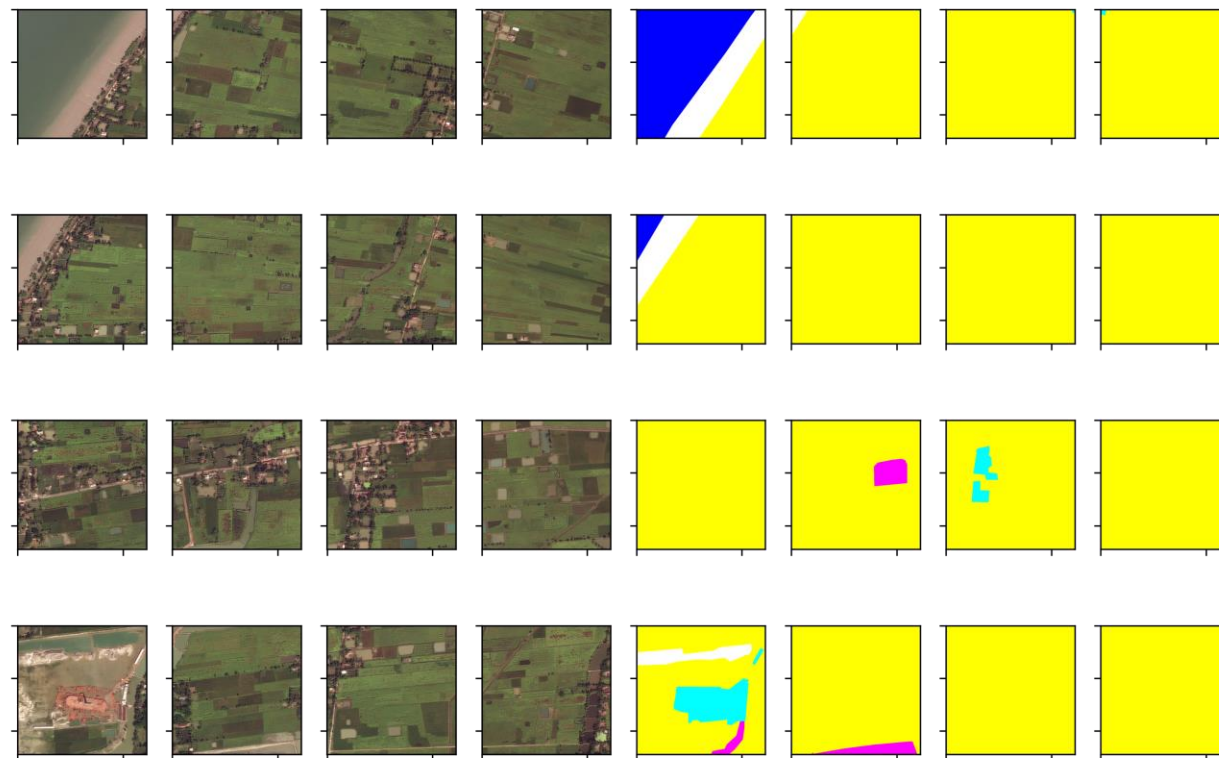
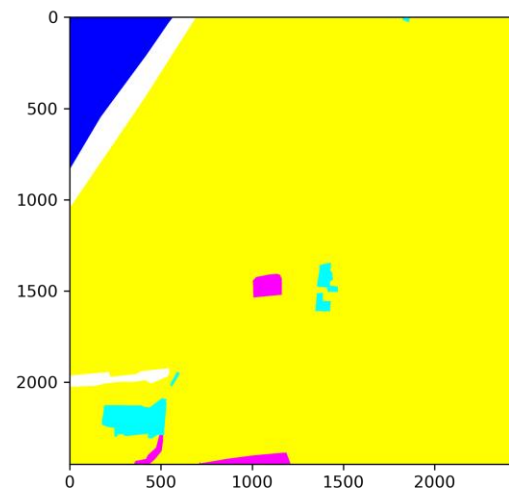

```
[55]: #copies all layers of model
activation_layers = [layer.output for layer in unet.layers]
# Creates a model that will return these outputs, given the model input
activation_model = models.Model(inputs=unet.input, outputs=activation_layers)
# Returns a list of five Numpy arrays: one array per layer activation
print("num of layers: ", len(activation_layers))
activation_model.summary()
```

num of layers: 72

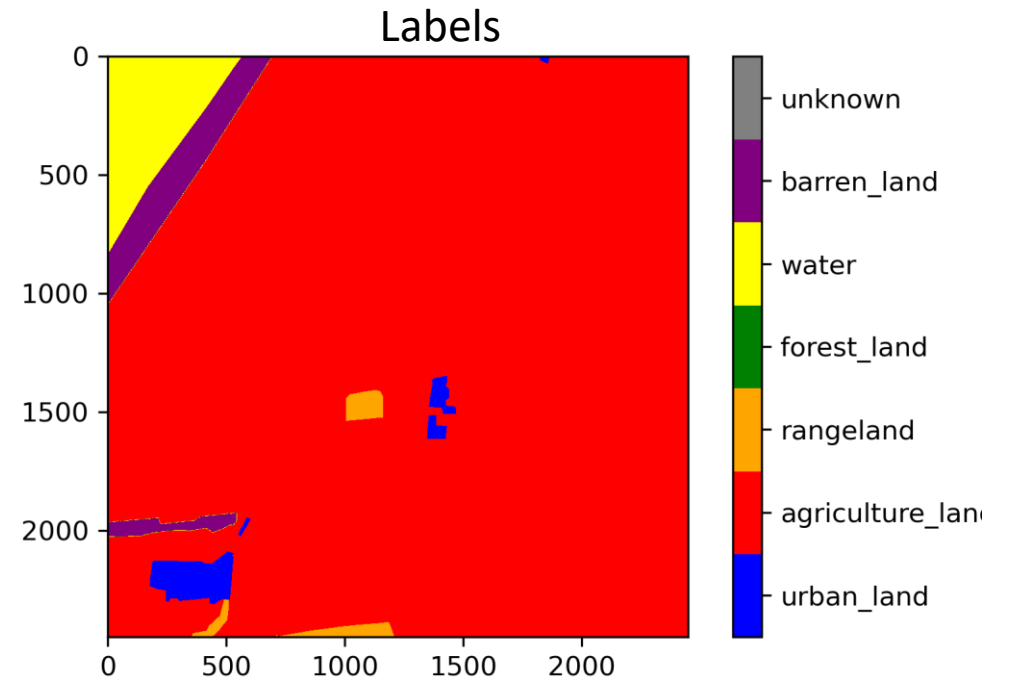
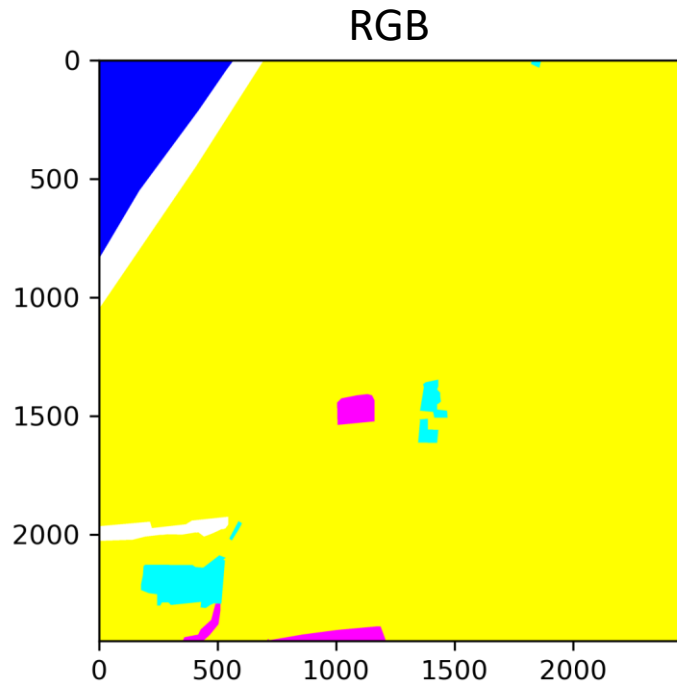
Model: "model_3"

Layer (type)	Output Shape	Param #	Connected to
=====			
Total params: 2,058,401			
Trainable params: 2,054,625			
Non-trainable params: 3,776			

Image Slicing



Convert Mask to Labels



Processing Satellite image

- Convert from jp2 with 13 bands to png with just 3 bands (RBG)




1st Dataset

- Image (64x64x3)
- Using ResNet



	Precision	Recall	F-Score	Support
AnnualCrop	0.950658	0.963333	0.956954	600.0
Forest	0.969055	0.991667	0.980231	600.0
HerbaceousVegetation	0.979522	0.956667	0.967960	600.0
Highway	0.979424	0.952000	0.965517	500.0
Industrial	0.975510	0.956000	0.965657	500.0
Pasture	0.972081	0.957500	0.964736	400.0
PermanentCrop	0.948104	0.950000	0.949051	500.0
Residential	0.946288	0.998333	0.971614	600.0
River	0.966601	0.984000	0.975223	500.0
SeaLake	0.998273	0.963333	0.980492	600.0

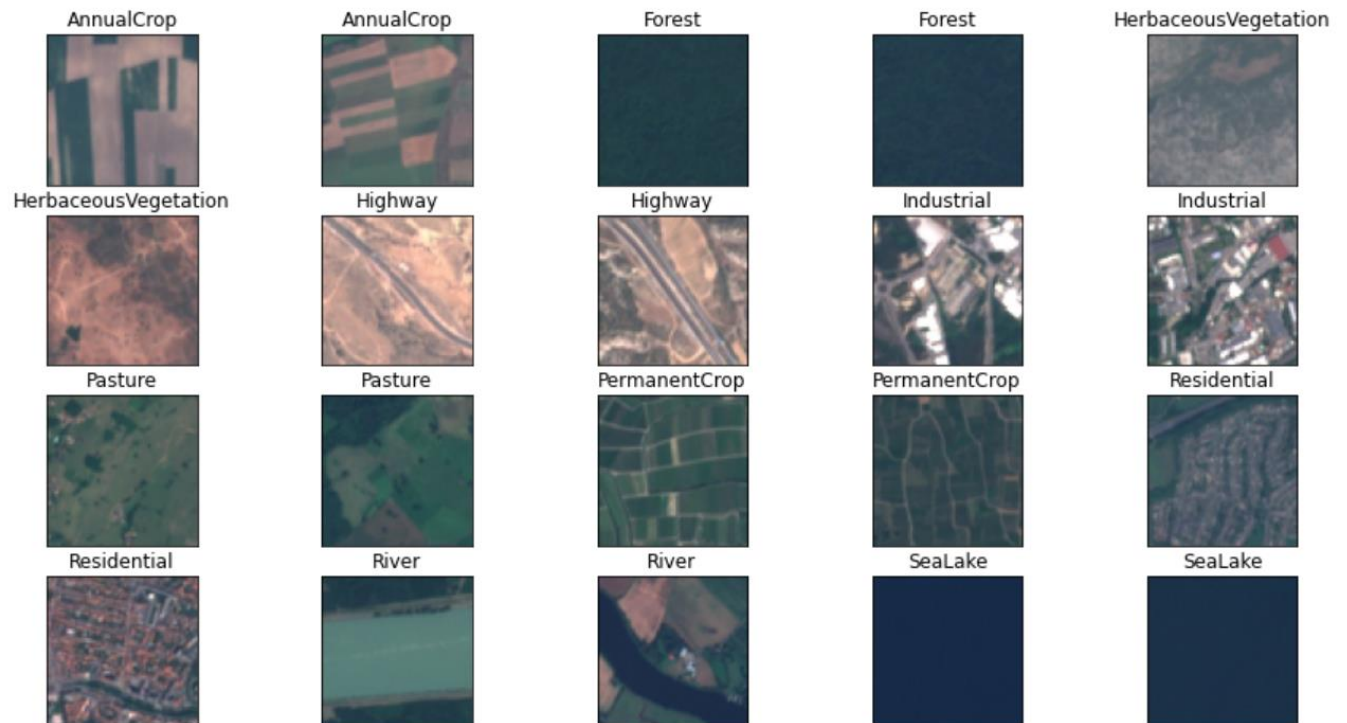


	AnnualCrop	Forest	HerbaceousVegetation	Highway	Industrial	Pasture	PermanentCrop	Residential	River	SeaLake
AnnualCrop	578	0	0	2	0	2	14	0	4	0
Forest	0	595	0	0	0	1	0	3	0	1
HerbaceousVegetation	3	5	574	0	1	5	9	3	0	0
Highway	3	0	1	476	8	2	1	3	6	0
Industrial	0	0	0	1	478	0	1	19	1	0
Pasture	7	4	5	0	0	383	1	0	0	0
PermanentCrop	9	0	6	1	2	0	475	6	1	0
Residential	0	0	0	0	1	0	0	599	0	0
River	2	0	0	6	0	0	0	0	492	0
SeaLake	6	10	0	0	0	1	0	0	5	578

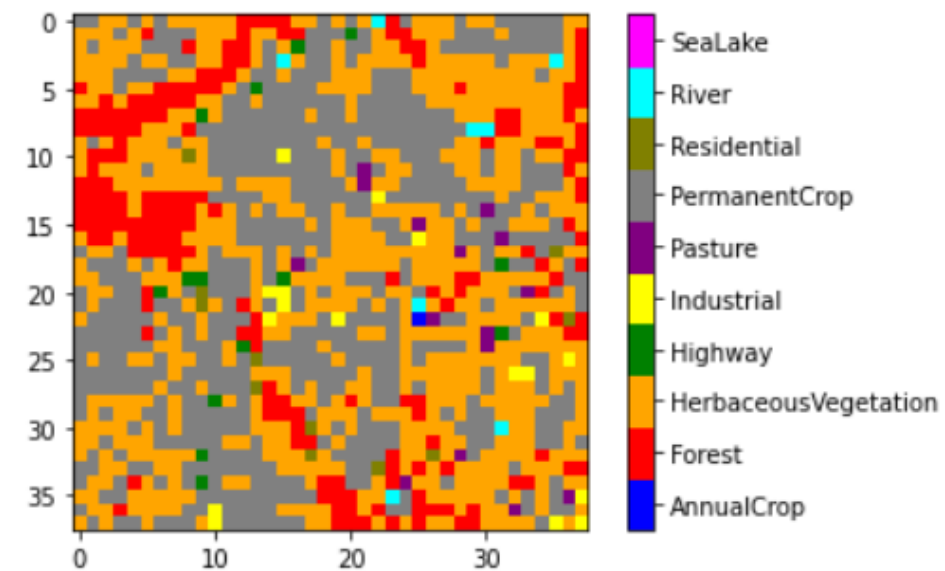
Results

- Image (64x64x3)
- Using ResNet

```
number_images = [1,500,616,1100,1230,1696,1820,2001,2400,2700,2902,3100,3300,3600,3800,4100,4500,4700,4900,5100]
plt.figure(figsize = (15, 8))
for i, r in enumerate(number_images):
    plt.subplot(4, 5, i+1, xticks = [], yticks = [])
    plt.imshow(test_generator_new[r][0]) # [the image from the images selected in the generator][there is just one image so 0 takes that image]
    plt.title(class_indices[predicted_classes[r]])
```




2448x2448x3

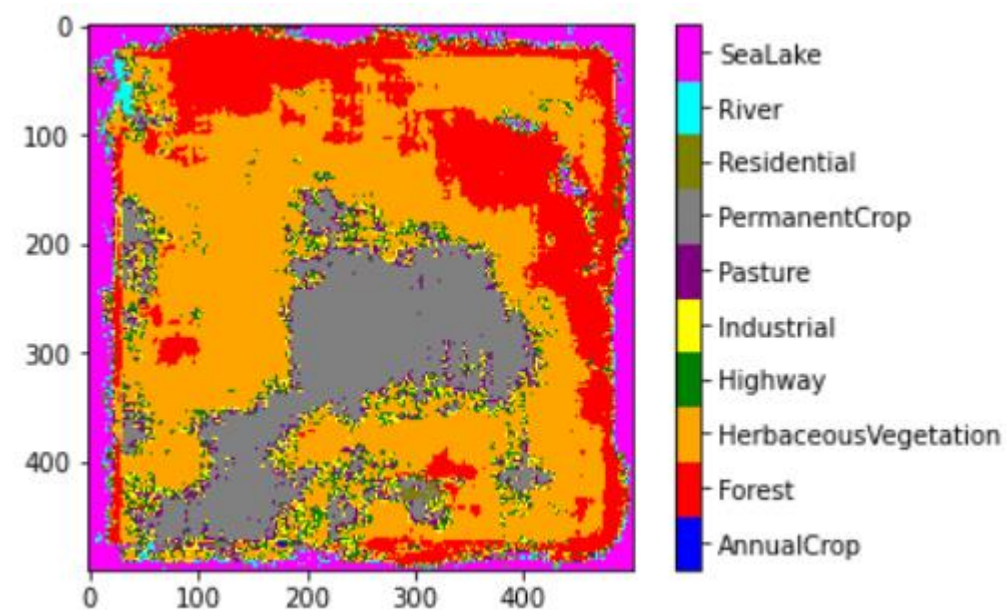


64x64

564x564x3



64x64



64x64 padded image

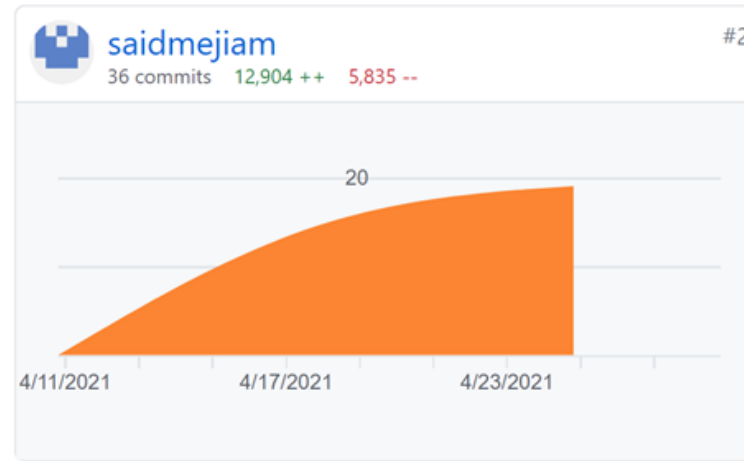

2nd Dataset

- Image (64x64x3)
- Using SegNet

```
def segnet(input_shape, n_labels, kernel=3, pool_size=(2, 2), output_mode="softmax"):  
    # encoder  
    inputs = Input(shape=input_shape)  
  
    conv_1 = Convolution2D(64, (kernel, kernel), padding="same")(inputs)  
    conv_1 = BatchNormalization()(conv_1)  
    conv_1 = Activation("relu")(conv_1)  
    conv_2 = Convolution2D(64, (kernel, kernel), padding="same")(conv_1)  
    conv_2 = BatchNormalization()(conv_2)  
    conv_2 = Activation("relu")(conv_2)  
  
    pool_1, mask_1 = MaxPoolingWithArgmax2D(pool_size)(conv_2)  
  
    conv_3 = Convolution2D(128, (kernel, kernel), padding="same")(pool_1)  
    conv_3 = BatchNormalization()(conv_3)  
    conv_3 = Activation("relu")(conv_3)  
    conv_4 = Convolution2D(128, (kernel, kernel), padding="same")(conv_3)  
    conv_4 = BatchNormalization()(conv_4)  
    conv_4 = Activation("relu")(conv_4)  
  
    pool_2, mask_2 = MaxPoolingWithArgmax2D(pool_size)(conv_4)  
  
    conv_5 = Convolution2D(256, (kernel, kernel), padding="same")(pool_2)  
    conv_5 = BatchNormalization()(conv_5)  
    conv_5 = Activation("relu")(conv_5)  
    conv_6 = Convolution2D(256, (kernel, kernel), padding="same")(conv_5)  
    conv_6 = BatchNormalization()(conv_6)  
    conv_6 = Activation("relu")(conv_6)  
    conv_7 = Convolution2D(256, (kernel, kernel), padding="same")(conv_6)  
    conv_7 = BatchNormalization()(conv_7)  
    conv_7 = Activation("relu")(conv_7)  
  
    pool_3, mask_3 = MaxPoolingWithArgmax2D(pool_size)(conv_7)  
  
    conv_8 = Convolution2D(512, (kernel, kernel), padding="same")(pool_3)  
    conv_8 = BatchNormalization()(conv_8)  
    conv_8 = Activation("relu")(conv_8)  
    conv_9 = Convolution2D(512, (kernel, kernel), padding="same")(conv_8)  
    conv_9 = BatchNormalization()(conv_9)  
    conv_9 = Activation("relu")(conv_9)  
    conv_10 = Convolution2D(512, (kernel, kernel), padding="same")(conv_9)  
    conv_10 = BatchNormalization()(conv_10)  
    conv_10 = Activation("relu")(conv_10)  
  
    pool_4, mask_4 = MaxPoolingWithArgmax2D(pool_size)(conv_10)  
  
    conv_11 = Convolution2D(512, (kernel, kernel), padding="same")(pool_4)  
    conv_11 = BatchNormalization()(conv_11)  
    conv_11 = Activation("relu")(conv_11)  
    conv_12 = Convolution2D(512, (kernel, kernel), padding="same")(conv_11)  
    conv_12 = BatchNormalization()(conv_12)  
    conv_12 = Activation("relu")(conv_12)  
    conv_13 = Convolution2D(512, (kernel, kernel), padding="same")(conv_12)  
    conv_13 = BatchNormalization()(conv_13)  
    conv_13 = Activation("relu")(conv_13)
```

```
# decoder  
unpool_1 = MaxUnpooling2D(pool_size)([pool_5, mask_5])  
  
conv_14 = Convolution2D(512, (kernel, kernel), padding="same")(unpool_1)  
conv_14 = BatchNormalization()(conv_14)  
conv_14 = Activation("relu")(conv_14)  
conv_15 = Convolution2D(512, (kernel, kernel), padding="same")(conv_14)  
conv_15 = BatchNormalization()(conv_15)  
conv_15 = Activation("relu")(conv_15)  
conv_16 = Convolution2D(512, (kernel, kernel), padding="same")(conv_15)  
conv_16 = BatchNormalization()(conv_16)  
conv_16 = Activation("relu")(conv_16)  
  
unpool_2 = MaxUnpooling2D(pool_size)([conv_16, mask_4])  
  
conv_17 = Convolution2D(512, (kernel, kernel), padding="same")(unpool_2)  
conv_17 = BatchNormalization()(conv_17)  
conv_17 = Activation("relu")(conv_17)  
conv_18 = Convolution2D(512, (kernel, kernel), padding="same")(conv_17)  
conv_18 = BatchNormalization()(conv_18)  
conv_18 = Activation("relu")(conv_18)  
conv_19 = Convolution2D(256, (kernel, kernel), padding="same")(conv_18)  
conv_19 = BatchNormalization()(conv_19)  
conv_19 = Activation("relu")(conv_19)  
  
unpool_3 = MaxUnpooling2D(pool_size)([conv_19, mask_3])  
  
conv_20 = Convolution2D(256, (kernel, kernel), padding="same")(unpool_3)  
conv_20 = BatchNormalization()(conv_20)  
conv_20 = Activation("relu")(conv_20)  
conv_21 = Convolution2D(256, (kernel, kernel), padding="same")(conv_20)  
conv_21 = BatchNormalization()(conv_21)  
conv_21 = Activation("relu")(conv_21)  
conv_22 = Convolution2D(128, (kernel, kernel), padding="same")(conv_21)  
conv_22 = BatchNormalization()(conv_22)  
conv_22 = Activation("relu")(conv_22)  
  
unpool_4 = MaxUnpooling2D(pool_size)([conv_22, mask_2])  
  
conv_23 = Convolution2D(128, (kernel, kernel), padding="same")(unpool_4)  
conv_23 = BatchNormalization()(conv_23)  
conv_23 = Activation("relu")(conv_23)  
conv_24 = Convolution2D(64, (kernel, kernel), padding="same")(conv_23)  
conv_24 = BatchNormalization()(conv_24)  
conv_24 = Activation("relu")(conv_24)  
  
unpool_5 = MaxUnpooling2D(pool_size)([conv_24, mask_1])  
  
conv_25 = Convolution2D(64, (kernel, kernel), padding="same")(unpool_5)  
conv_25 = BatchNormalization()(conv_25)  
conv_25 = Activation("relu")(conv_25)  
  
conv_26 = Convolution2D(n_labels, (1, 1), padding="valid")(conv_25)  
conv_26 = BatchNormalization()(conv_26)  
conv_26 = Reshape(  
    (input_shape[0] * input_shape[1], n_labels),  
    input_shape=(input_shape[0], input_shape[1], n_labels),  
) (conv_26)  
  
outputs = Activation(output_mode)(conv_26)  
# print("Build decoder done..")  
  
model = Model(inputs=inputs, outputs=outputs, name="SegNet")  
  
return model
```


Git hub commits 151 total



Notebook total

- HG: 4
- JB: 7
- JM: 4
- SM: 3